



Module R3/R4

Dispatching

Introduction

- Modules R3 and R4 are due at the same time
- Implement context switching and support for multiple processes
- Majority of the code will be written in C, but some Assembly is required (lol)

R3 Goals

- Track particular processes with PCB's
- Support `sys_req(IDLE, ...)` and `sys_req(EXIT, ...)`
- Switch out the running process
- Add temporary commands to test above
 - Adding new processes
 - Execute processes until all are complete, then return to `comm_hand`

R4 Goals

- Treat comm_hand as just another process
- Prevent your polling function from blocking
- Remove temporary commands
- Add an alarm command
- Add an infinitely running process, that must be user suspended to be killed

Review: A Program in Execution

- The CPU has assorted registers
 - Most notably, the Instruction Pointer, the Stack Pointer, and the Base (of the Stack) Pointer
- All of the variables are stored on the stack
 - Heaps are done in R5
- That's all the values necessary to save, load, fork a running process, and are collectively referred to as the 'context'

How Interrupts Work

- Similar to a function call
 - A 'return address' is pushed to the stack
 - Jump to the interrupt table
 - Execute the function at that spot in the table
 - Jump back to the return address
- The interrupt is responsible for preserving the register values
 - If the interrupt needs to use/change a register value, it must push it at the beginning of the function and pop it at the end of the function

Context Switch

- What is a context switch?
 - An event where one process gives up the processor for another process.

Overview of a context switch:

 - ▢ Save the hardware state (i.e. the values of the CPU registers) of the currently-operating process
 - ▢ Reinsert the currently-operating process into the appropriate queue
 - ▢ Select the next process to begin or continue executing
- What is a context?
 - The top 60 bytes of the stack
 - Contains the 15 CPU register values needed to begin and resume execution

```
struct context {  
    u32int gs, fs, es, ds;  
    u32int edi, esi, ebp, esp, ebx, edx, ecx, eax;  
    u32int eip, cs, eflags;  
};
```
- Look here for more info
 - <http://int.cs.yale.edu/cs422/doc/pc-arch.html>

Context Switching

- How does the system decide when a context switch will occur?
 - In MPX, a process voluntarily gives up control of the CPU by generating a software interrupt
 - Processes will call `sys_req(IDLE)`; to give up control
 - `sys_req(IDLE)`; will generate an interrupt on line 60h
 - ▢ Check out `mpx_supt.h` for more info about `sys_req`
 - Your interrupt handler will perform the context switch.

Prerequisite Knowledge

- What is the calling convention gcc uses for IA-32 CPUs?
 - Function parameters are pushed onto the stack in reverse-order (that is: the first parameter is the last to be pushed onto the stack)
 - The lower 32 bits of the return value (all we care about for this project) is stored in the EAX register
- What must an interrupt handler do when an interrupt occurs to avoid disturbing the state of a running process?
 - Save the contents of all general purpose registers (i.e. push the values onto the stack)
 - Service the interrupt
 - Restore the computer's state (all CPU register values)

Prerequisite Knowledge

- How can the aforementioned description of an interrupt handler be used to perform a context switch?
 - Simple! - Dedicate a stack to each process and "switch" stacks. For example:
 - Push all general purpose CPU register values onto the stack
 - Switch stacks
 - Pop all general purpose CPU register values off of the stack - Note that this pops different values into the registers than those saved in step 1.
- How much assembly language do I need to know?
 - pusha - Push all general purpose CPU register values onto the stack
 - popa - Pop values off of the stack into all general purpose CPU registers
 - push - Push the value given by the operand onto the stack
 - pop - Pop the value off of the top of the stack into the location described by the operand
 - mov - Move the contents of one location into another
 - call - Call a subroutine
 - iret - Return from interrupt

R3

- For R3 and R4, you will need to add two new commands (more on this later), fill in the body of the system call interrupt service routine (in `irq.s`) and write a system call function which prepares MPX for the next ready process to begin/resume execution.
- You may write your system call function in C or assembly
 - C is advised
 - `u32int* sys call(context *registers)`
- GCC is incapable of creating "naked functions", so the Interrupt Service Routine stub must be written in assembly language

sys_call_isr

- Push all general purpose register values to the stack
- Push the segment register values to the stack (in the order: ds, es, fs, gs)
- Push esp, a register serving as an indirect memory operand pointing to the top of the stack at any time, this will be the parameter of your sys_call function
- Call sys_call
- Set a new stack pointer (returned by your sys_call function)
 - Recall that the return value of your sys_call function will be stored in eax
- Pop values into your segment registers in reverse order (from step 2)
- Pop values into general purpose registers
- Return from interrupt
- This is one of the easiest parts of R3/R4, come to me if you find yourself struggling, I don't want you to be stuck here

sys_call

- Declare a PCB* cop as a global variable, representing the currently operating process
- Use the prototype `u32int* sys_call(context* registers)`
- Check to see if sys call has been called before (i.e. if the currently-operating process (cop) is NULL, sys call has not been called).
 - If sys_call has not been called, save a reference to old (the caller's) context in a global variable.
 - If sys_call has been called check params.op code
 - ▢ If params.op code == IDLE, save the context (reassign cop's stack top)
 - ▢ If params.op code == EXIT, free cop
- If there is a ready process:
 - Remove it from the ready queue, set state to running
 - Assign cop
 - Return cop's stack top
- Otherwise, return the context saved in step 1

User Commands

- yield
 - The yield command (temporary – only in R3) will cause command to yield to other processes (i.e. voluntarily give up CPU time)
 - If any processes are in the ready queue, they will be executed
 - Yield is literally a one-liner:
 - `asm volatile ("int $60");`
- loadr3
 - loadr3 will load all r3 "processes" (proc3.c file eCampus) into memory in a suspended ready state at any priority of your choosing

```
pcb * new_pcb = create_pcb ( name , 1 , 1 , 1 , stack_size );
context * cp = ( context * )( new_pcb -> stack_top );
memset ( cp , 0, sizeof ( context ));
cp -> fs = 0 x10 ;
cp -> gs = 0 x10 ;
cp -> ds = 0 x10 ;
cp -> es = 0 x10 ;
cp -> cs = 0 x8 ;
cp -> ebp = ( u32int )( new_pcb -> stack );
cp -> esp = ( u32int )( new_pcb -> stack_top );
cp -> eip = ( u32int ) func ;// The function correlating to the process , ie. Proc1
cp -> eflags = 0 x202 ;
return new_pcb ;
```

VERY IMPORTANT

- In order for your interrupt handler to hook an interrupt request on line 60, you need to enter the appropriate interrupt handler in the interrupt descriptor table
(kernel/core/interrupts.c)
 - `idt_set_gate(60, (u32int)sys_call_isr, 0x08, 0x8e);`
- In `setupPCB()` from R2, be sure to set your stacktop to:
 - `stackbase + 1024 - sizeof(struct context)`

R4

- R4 is to demonstrate continuous dispatching and to allow commhand to compete for resources (it becomes a process)
- The deletePCB() command is no longer temporary due to new changes, and will be used in this module
- Inside kmain(), the following will need to happen:
 - Remove the call to commhand
 - Add commhand and idle to your ready queues (see the previous slide concerning loading R3 processes).
 - Trigger a software interrupt on line 60
- Remove the yield command from the user

R4 - Alarm

- Create an alarm command/process with 2 components
 - Message – Some user specified string
 - Time – When message should be printed
 - Process idles when a time has not yet occurred
 - Process prints the message and exits at or after that time
- System should support multiple concurrent alarms
- It would be wise to keep the alarm PCB parameterless, and upon its creation/deletion, update a list of times for it to check

R4 - Infinite Command/Process

- Should forever call `sys_req(IDLE...)`
 - If not suspended, the process can not be deleted
 - If suspended, it may be deleted
- Should be implemented as a process
 - Process idles when not yet suspended
 - Process prints a message every time it runs, to show it is properly running



Questions?